



THE UNIVERSITY *of* EDINBURGH



AI accelerators for HPC

Joseph Lee

EPCC

j.lee@epcc.ed.ac.uk

Who are we?


- Supercomputing centre at the University of Edinburgh, UK
- Providing world-class computing systems, data storage and support services
- Hosts the UK's national supercomputing service: ARCHER2



- ~150 staff
- Research, Education and Training: MSc and PhD programme

Who are we?

- Edinburgh International Data Facility (EIDF)
- High-powered data analytics and storage service that supports research and data-driven innovation
 - CPU & GPU
 - Cerebras CS-2 (x2)
 - Graphcore Bow Pod 64
- We have also been using these machines for HPC research



What are we looking at?

- ExCALIBUR (UK research programme for exascale algorithms and infrastructure)
- Enabling Hardware & Software Programme
 - Accelerators: RISC-V, FPGAs, Cerebras WSE



TensorFlow as a DSL for stencil-based
computation on the Cerebras Wafer Scale Engine

Nick Brown¹, Brandon Echols², Justs Zarins¹, and Tobias Grosser³

- xDSL (Python x MLIR) for stencil computation
- PETSc on Cerebras
 - And more we want to do!
- Training: SDK tutorial at HPC Days – Durham UK



Experience

Positive 👍	Challenges 😞
Helpful support	Programming model/language (DSL)
Active community	Training & documentation
Tools (e.g. simulator)	Closed tools & reproducibility
Performance!	Big software changes
	Repeated/similar work (more collaboration?)

What are we looking for?

- Usable and accessible: Language, Library, Compiler
- Documentation and support
- AI Accelerator for HPC tutorials?
- Community!